

Major Exam (COL 351)

Read the instructions carefully:

- You need to justify correctness and running time of each algorithm.
- If using dynamic programming, you must explain the meaning of table entries, and explain the order of computing them.
- For NP-completeness reduction, you can use the fact that 3-SAT, Independent Set, Vertex Cover, Partition, Subset Sum, Clique and Hamiltonian Cycle (in directed and in undirected graphs) are NP-complete.
- No argument should use examples – they will be ignored.
- You can assume any result proved during the lectures (but cannot assume any other result which is in the book or in the tutorials).

1. You are organizing a sports event, where each contestant has to swim 20 rounds in a pool, and then run 3 kilometers. However, the pool can be used by only one person at a time. In other words, the first contestant swims 20 rounds, gets out and then starts running. As soon as this first person is out of the pool, a second contestant begins swimming the 20 rounds; as soon as he/she is out and starts running, a third contestant begins swimming... and so on.) (6 marks)

You are given a list of n contestants, and for each contestant you are given the time it will take him/her to complete swimming 20 rounds of the pool, and the time it will take him/her to run 3 kilometers. Your job is to decide on a schedule for the event, i.e., an order in which to sequence the starts of the contestants. The completion time of a schedule is the earliest time at which all contestants will be finished with swimming and running. Give an efficient algorithm that produces a schedule whose completion time is as small as possible.

Example: Suppose there are two contestants C_1, C_2 with swimming and running times being (10, 5) for C_1 and (2, 8) for C_2 . If we schedule them as C_1, C_2 , then C_1 will finish swimming and running by time $10 + 5 = 15$. C_2 can start swimming at time 10, and so, will finish by time $10 + 2 + 8 = 20$. Note that both contestants will finish by time 20, and so the completion time is 20. If we order them as C_2, C_1 , then C_2 will finish by time 10 and C_1 by time 17. Therefore, the completion time is 17. Thus, the best ordering is C_2, C_1 .

Solution: The algorithm schedules the contestants in decreasing order of their running time (2 marks). Let this order be C_1, \dots, C_n . Now we prove correctness. Consider an optimal schedule O_1, \dots, O_n . If it is same as C_1, \dots, C_n , then we have nothing to prove. So assume this is not the case. Let R_i denote the running time and S_i denote the swimming time of contestant O_i . Let T denote the time by which all contestants finish. Then there will be two consecutive contestants O_i, O_{i+1} such that $R_i \leq R_{i+1}$ (1 mark). Let T_i be the time at which O_i starts swimming. Then O_i finishes at time $T_i + S_i + R_i$, whereas O_{i+1} finishes at time $T_i + S_i + S_{i+1} + R_{i+1}$. So,

$$T \geq T_i + S_i + S_{i+1} + R_{i+1}.$$

Now we consider a new schedule which is same as O but swaps O_i and O_{i+1} (1 mark). Here, the finishing time of O_i and O_{i+1} are $T_i + S_i + S_{i+1} + R_i$ and $T_i + S_{i+1} + R_{i+1}$, and the other contestants' finishing time remains unchanged. Since $R_i \leq R_{i+1}$, we see that $T_i + S_i + S_{i+1} + R_i, T_i + S_{i+1} + R_{i+1} \leq T_i + S_i + S_{i+1} + R_{i+1} \leq T$. Therefore, the new schedule is not worse than O (1 mark). By continuing

this swapping process, we will get the greedy schedule, and the completion time would be no worse than T (**1 mark**).

Common Mistakes: Many students have got 5 out of 6, because they did not mention the last step. You need to continue such swaps to get to greedy. Many of them assumed that $R_i < R_{i+1}$, and so, one would get a strict improvement. But R_i could be equal to R_{i+1} .

In case you got the greedy rule wrong, I have given 0 or 1 marks depending on whether the approach for proof looked ok. This is more of a subjective assessment and I will not entertain any queries on this.

- Let S be a set of $n + 1$ distinct integers. You can assume that S is given as an array. You are given an unsorted array A of size n containing exactly n out of the $n + 1$ integers in S . Give an $O(n)$ time algorithm to find the integer from S which is not in A . The only operation allowed on numbers in S (or A) is comparison (you are NOT allowed to perform addition, subtraction, multiplication, etc. on these numbers). (**5 marks**)

Solution: We use divide and conquer strategy. First find the median of A in $O(n)$ time (**2 mark**). Call this x . Let A_L be the numbers in A which are less than or equal to x and S_L be the numbers less than or equal to x in S . Define A_R and S_R similarly (**1 mark**). Now, either $|A_L| = |S_L| - 1$, or $|A_R| = |S_R| - 1$. Depending on the case, we proceed recursively (**1 mark**). Since finding median takes $O(n)$ time, we get the recurrence $T(n) = T(n/2) + O(n)$ (**1 mark**), whose solution is $T(n) = O(n)$. If you used randomized version (like quick-sort), that is ok.

- You are given an array A containing n numbers (which could be positive, zero or negative rational numbers). A sub-array $A[i, j]$ of A , where $i \leq j$, is defined by the sequence $A[i], A[i + 1], \dots, A[j]$. For each such sub-array, define $P(i, j)$ as the product of the entries in $A[i, j]$. Give an $O(n)$ time algorithm to find the largest value of $P(i, j)$ overall sub-arrays $A[i, j]$ (note that the algorithm just outputs a number). You can assume that arithmetic operations like multiplication on numbers in A take constant time. (**6 marks**)

Example: Suppose A is $\{-3, 10, -6, 7, 2, -1\}$. Assuming that the first element of A is denoted by $A[1]$, the sub-array $A[1, 4]$ has total product $-3 \times 10 \times -6 \times 7 = 1260$, whereas sub-array $A[3, 6]$ has total product $-6 \times 7 \times 2 \times -1 = 84$.

Solution: We use dynamic programming. We maintain two tables $S[i]$ and $L[i]$, for $i = 1, \dots, n$. The entry $L[i]$ denotes the largest value of $A_i \cdot A_{i+1} \cdots A_j$ for $j \geq i$, and similarly, $S[i]$ denotes the smallest value of $A_i \cdot A_{i+1} \cdots A_j$ for all $j \geq i$ (**2 marks**). Base case is easy: $L[n]$ and $S[n]$ are both equal to A_n (**0.5 mark**). Now, for computing $L[i]$, we proceed depending on A_i being positive or negative. Assume the first case. Then observe that $L[i] = \max(A_i, A_i \cdot L[i + 1])$, and $S[i] = \min(A_i, A_i \cdot S[i + 1])$ (**1 mark**). If A_i were negative, then $L[i] = \max(A_i, A_i \cdot S[i + 1])$, and $S[i] = \min(A_i, A_i \cdot L[i + 1])$ (**1.5 mark**). Finally, we output $\max_i L_i$ (**0.5 mark**). We compute the entries starting from $i = n$ and decreasing i (**0.5 mark**).

Many people got the definition of $L[i]$ (or $S[i]$ wrong) – they just said the maximum produce in $A_i \dots A_n$. But this is wrong – it needs to include A_i . Many of them define $S[i]$ as the largest negative number. This is wrong, but then one has to be careful with initial conditions – for example, what if all entries so far are positive? I have taken off 1 mark if one is not careful with the initial condition.

- A town has r residents R_1, \dots, R_r , q clubs C_1, \dots, C_q and p political parties P_1, \dots, P_p . Each resident is a member of exactly one club and belongs to exactly one political party. You want to form a governing council for the town. The governing council must contain exactly l_i members from club C_i , for $i = 1, \dots, q$; but it can have at most u_k members from the political party P_k , for $k = 1, \dots, p$. Give

an efficient algorithm which either finds such a council or declares that no such council is possible. Assume that the membership information is given in a suitable data-structure. **(4 marks)**

Solution: We reduce this problem to max-flow. We have a node c_i for each club C_i , a node p_i for each political party P_i and a node r_i for each resident R_i . If R_i is a member of club C_j and party P_l , then we have an edge from c_j to r_i and another from r_i to p_l , both of capacity 1. We add a source node s and a sink node t . We have edges from s to c_i with capacity l_i and from p_i to t of capacity u_i . We now find a max-flow from s to t , and check if its value is equal to $\sum_i l_i$.

To prove correctness, note that if there is a solution, then we can send a flow as follows: we send l_i units of flow from s to c_i , and 1 unit of flow from c_i to every r_j for those residents of C_i who get chosen. If R_i is selected, we send 1 unit of flow from r_i to the corresponding political party node, and then from every p_i the flow from p_i to t is equal to the flow received by p_i . Conversely, if there is a flow saturating all the edges (s, c_i) , then by integrality of flow, each r_j receives either 0 or 1 unit of flow. If r_j receives 1 unit of flow, then we select R_i . Again, it is easy to check that this satisfies all the constraints.

Common Mistakes: There has been binary marking here – you either got this right or not. Many people missed out saying that given a valid flow, one needs integrality property to recover a solution. I have taken off 0.5 marks for this. Similarly many people mentioned wrong capacities – I have taken off 1 or 2 marks. If you got the graph wrong, you got maximum of 1 mark – again, a subjective assessment depending on your solution.

5. Consider the following optimization version of PARTITION problem. You are given n integers x_1, \dots, x_n and would like to partition them into sets A and B such that $\max(\text{Sum}(A), \text{Sum}(B))$ is minimized, where $\text{Sum}(X)$ denotes the sum of all the numbers in X . Given a solution A, B , define its value as the quantity $\max(\text{Sum}(A), \text{Sum}(B))$. Consider the following greedy algorithm – initialize sets A and B to emptyset. Consider the numbers x_1, \dots, x_n iteratively. When looking at x_i , if $\text{Sum}(A) < \text{Sum}(B)$, add x_i it to A , else add it to B .

- (a) Prove that there is an example for which this algorithm has value $3/2$ times the minimum possible value for this example. **(1 mark)**

Solution: Consider the numbers 1,1,2.

- (b) Prove that this algorithm has the property that for any input, its value is at most $3/2$ times the minimum possible value for this input. **(6 marks)**

Solution: Consider an input x_1, \dots, x_n . let A, B denote the sets constructed by our algorithm, and let O be the optimal value for this input. Assume that $\text{Sum}(A) \geq \text{Sum}(B)$, and let Σ denote the sum of all the numbers x_1, \dots, x_n . Let x_i be the last number added to A , and suppose the sum of numbers in A just before adding x_i was S **(1 mark)**. Then our algorithm has value $x_i + S$ **(1 mark)**. Also, the sum of B would be at least S (otherwise we would not add x_i to A **(1 mark)**). So $\Sigma \geq 2S + x_i$. Now observe that $O \geq x_i$ **(1 mark)** and $O \geq \Sigma/2 \geq S + x_i/2$ **(1 mark)**. Combining these, we get $3O/2 \geq x_i + S$, which is the value of our solution **(1 mark)**.

- (c) Suggest an algorithm which would have better ratio than $3/2$ – you need not prove anything, just give a one line algorithm. **(1 mark)**

Solution: Run the same algorithm as above, but first sort x_i in decreasing order. If say “sorted order”, you get 0.5 marks. If you say increasing order, it is clearly wrong.

6. Consider the following variant of max-flow: you are given a directed graph G with positive integer edge capacities, two vertices s and t . You would like to find a flow from s to t such that flow on every edge e is either 0 or u_e , where u_e denotes the capacity of e – we call such a flow a “saturating flow”.

Prove that the following problem is NP-complete: given a directed graph G with vertices s, t and positive integer edge capacities, and a parameter k , is there a saturating flow from s to t of value at least k ? (6 marks)

Solution: This problem is clearly in NP. A solution just needs to give a flow of value k (0.5 mark), and a verifier needs to check that it is a valid flow and saturates every edge with non-zero flow on it (0.5 mark). The most common mistake is that many people did not mention that one needs to check flow conservation constraint. A solution will only mention f_e values on each edge, so the verifier needs to check this. Most of you got 0.5 marks here.

We prove NP-completeness by reduction from the subset sum problem. Consider an instance I of the subset sum problem, where are given numbers x_1, x_2, \dots, x_n and a number S . We would like to answer if there is a subset of these numbers which add up to S . We construct an instance of the saturating flow problem as follows. We have a vertices s and s' . For every element x_i , we add a new vertex v_i , edges from s to v_i and from v_i to s' , both of capacity x_i . Now we add a new vertex t and add an edge from s' to t of capacity S . Now we claim that there is a solution to the subset sum problem if and only if there is a saturating flow from s to t of value S . To show this, suppose there is a solution to the subset sum problem, and let the subset be X . Then we can have a flow as follows – we send x_i unit of flow from s to v_i to s' if $x_i \in X$. Clearly the flow on s' to t would be S . Conversely, if there is a such a flow, then let x_i be the elements such that we have x_i flow on s to v_i edge. Then these elements add to S .

The most common mistake here is that students did not add the edge s' to t . Although this would have been correct if the problem definition had asked for a saturated flow of value exactly k , but it is not correct if one wants flow of value at least k . In other words, if you did not add this edge, a saturated flow of value at least S does not guarantee a solution to subset sum (which requires the sum to be exactly S). I have taken off 1.5 marks for this mistake.

Again the grading here has been binary – I could not find any correct reduction from a problem other than subset sum (or partition). Sometimes I gave 1 mark depending on how close you were to a correct reduction. Again, please don't come to get this extra mark.